

# Project Name : PREDICTIVE CUSTOMER BEHAVIOUR MODELLING USING AI



Predictive Customer Behavior Modelling Using AI is an in-depth analysis of a company's ideal customers. This approach enables businesses to gain a clearer understanding of their audience, making it easier to adapt products to meet the unique needs, preferences, and concerns of various customer groups. By leveraging predictive behavior modelling, organizations can tailor their offerings based on distinct customer segments. For example, instead of investing resources to promote a new product to every individual in the company's database, the business can identify which segment is most likely to purchase the product and focus its marketing efforts on that specific group.

## Time Line of the project :

- Importing Libraries
- Data Analysis
- Data Cleaning and Feature Engineering
- Performing Clustering

### Importing Libraries

```
import pandas as pd ## analysis
import numpy as np ## computational ability
import seaborn as sns
```

```

import matplotlib.pyplot as plt ## visualization
%matplotlib inline

df= pd.read_csv("Customer_Analytics_DB.csv",sep='\t')

df

```

	ID	Year_Birth	Education	Marital_Status	Income	
Kidhome \	5524	1957	Graduation	Single	58138.0	0
0	2174	1954	Graduation	Single	46344.0	1
1	4141	1965	Graduation	Together	71613.0	0
2	6182	1984	Graduation	Together	26646.0	1
3	5324	1981	PhD	Married	58293.0	1
4	...	...	...	...	...	...
...	...	...	...	...	...	...
2235	10870	1967	Graduation	Married	61223.0	0
2236	4001	1946	PhD	Together	64014.0	2
2237	7270	1981	Graduation	Divorced	56981.0	0
2238	8235	1956	Master	Together	69245.0	0
2239	9405	1954	PhD	Married	52869.0	1
Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth	
\	0	04-09-2012	58	635	...	7
0	1	08-03-2014	38	11	...	5
1	0	21-08-2013	26	426	...	4
2	0	10-02-2014	26	11	...	6
3	0	19-01-2014	94	173	...	5
4	...	...	...	...	...	...
...	...	...	...	...	...	...
2235	1	13-06-2013	46	709	...	5
2236	1	10-06-2014	56	406	...	7
2237	0	25-01-2014	91	908	...	6

2238	1	24-01-2014	8	428	...	3
2239	1	15-10-2012	40	84	...	7
AcceptedCmp2	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1		
0	0	0	0	0		
0						
1	0	0	0	0	0	
0						
2	0	0	0	0	0	
0						
3	0	0	0	0	0	
0						
4	0	0	0	0	0	
0						
...	...	...	...	...	...	
...						
2235	0	0	0	0	0	
0						
2236	0	0	0	0	1	
0						
2237	0	1	0	0	0	
0						
2238	0	0	0	0	0	
0						
2239	0	0	0	0	0	
0						
Complain	Z_CostContact	Z_Revenue	Response			
0	0	3	11	1		
1	0	3	11	0		
2	0	3	11	0		
3	0	3	11	0		
4	0	3	11	0		
...	...	...	...	...	...	
2235	0	3	11	0		
2236	0	3	11	0		
2237	0	3	11	0		
2238	0	3	11	0		
2239	0	3	11	1		

[2240 rows x 29 columns]

```
df.shape
(2240, 29)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   ID                2240 non-null    int64  
 1   Year_Birth        2240 non-null    int64  
 2   Education         2240 non-null    object  
 3   Marital_Status    2240 non-null    object  
 4   Income             2216 non-null    float64 
 5   Kidhome           2240 non-null    int64  
 6   Teenhome          2240 non-null    int64  
 7   Dt_Customer       2240 non-null    object  
 8   Recency            2240 non-null    int64  
 9   MntWines           2240 non-null    int64  
 10  MntFruits          2240 non-null    int64  
 11  MntMeatProducts   2240 non-null    int64  
 12  MntFishProducts   2240 non-null    int64  
 13  MntSweetProducts  2240 non-null    int64  
 14  MntGoldProds      2240 non-null    int64  
 15  NumDealsPurchases 2240 non-null    int64  
 16  NumWebPurchases   2240 non-null    int64  
 17  NumCatalogPurchases 2240 non-null    int64  
 18  NumStorePurchases 2240 non-null    int64  
 19  NumWebVisitsMonth 2240 non-null    int64  
 20  AcceptedCmp3      2240 non-null    int64  
 21  AcceptedCmp4      2240 non-null    int64  
 22  AcceptedCmp5      2240 non-null    int64  
 23  AcceptedCmp1      2240 non-null    int64  
 24  AcceptedCmp2      2240 non-null    int64  
 25  Complain           2240 non-null    int64  
 26  Z_CostContact     2240 non-null    int64  
 27  Z_Revenue           2240 non-null    int64  
 28  Response            2240 non-null    int64  
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB

```

df.dtypes

ID	int64
Year_Birth	int64
Education	object
Marital_Status	object
Income	float64
Kidhome	int64
Teenhome	int64
Dt_Customer	object
Recency	int64
MntWines	int64
MntFruits	int64

```

MntMeatProducts           int64
MntFishProducts           int64
MntSweetProducts          int64
MntGoldProds              int64
NumDealsPurchases         int64
NumWebPurchases           int64
NumCatalogPurchases       int64
NumStorePurchases         int64
NumWebVisitsMonth         int64
AcceptedCmp3               int64
AcceptedCmp4               int64
AcceptedCmp5               int64
AcceptedCmp1               int64
AcceptedCmp2               int64
Complain                  int64
Z_CostContact             int64
Z_Revenue                 int64
Response                  int64
dtype: object

```

## Data Analysis

Null Values

```

df.isnull().sum()

ID                      0
Year_Birth                0
Education                 0
Marital_Status             0
Income                     24
Kidhome                   0
Teenhome                   0
Dt_Customer                0
Recency                    0
MntWines                   0
MntFruits                  0
MntMeatProducts             0
MntFishProducts             0
MntSweetProducts            0
MntGoldProds                0
NumDealsPurchases           0
NumWebPurchases             0
NumCatalogPurchases          0
NumStorePurchases            0
NumWebVisitsMonth            0
AcceptedCmp3                 0
AcceptedCmp4                 0
AcceptedCmp5                 0

```

```
AcceptedCmp1          0
AcceptedCmp2          0
Complain              0
Z_CostContact         0
Z_Revenue              0
Response               0
dtype: int64

mean= df[ 'Income' ].mean()
df[ 'Income' ]= df[ 'Income' ].fillna(mean)
```

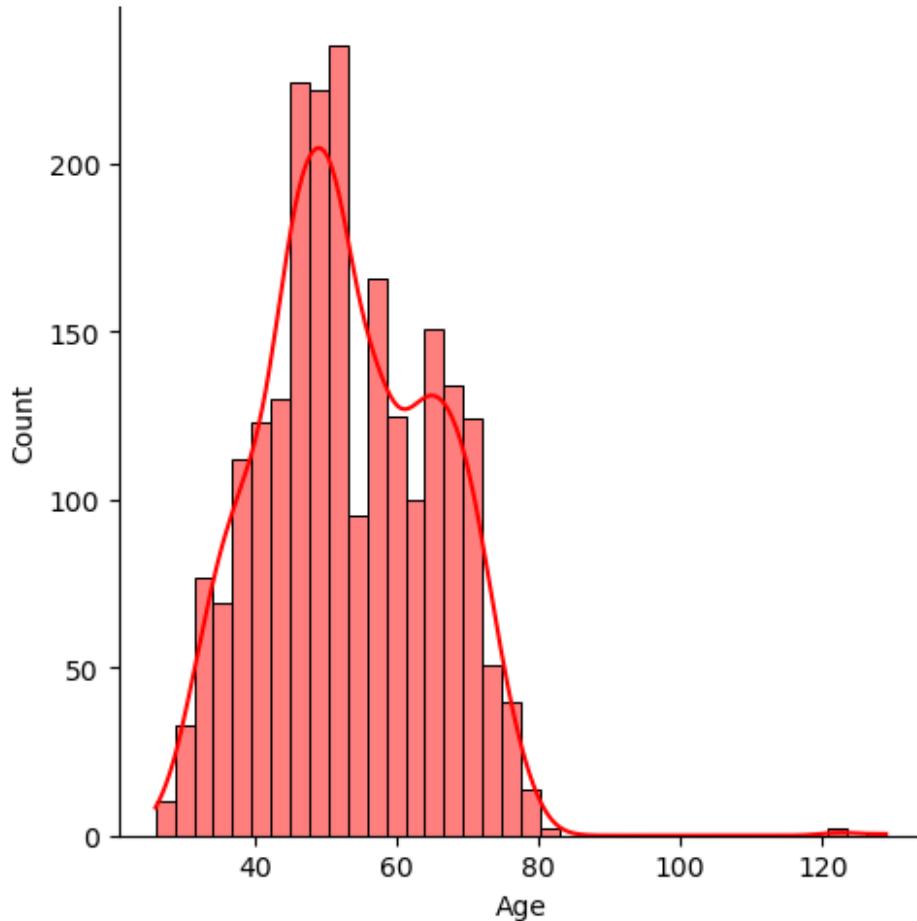
Let us create a new column Age by using year of birth

```
df[ 'Age' ]= 2022-df[ 'Year_Birth' ]

df[ 'Age' ]

0      65
1      68
2      57
3      38
4      41
 ..
2235    55
2236    76
2237    41
2238    66
2239    68
Name: Age, Length: 2240, dtype: int64

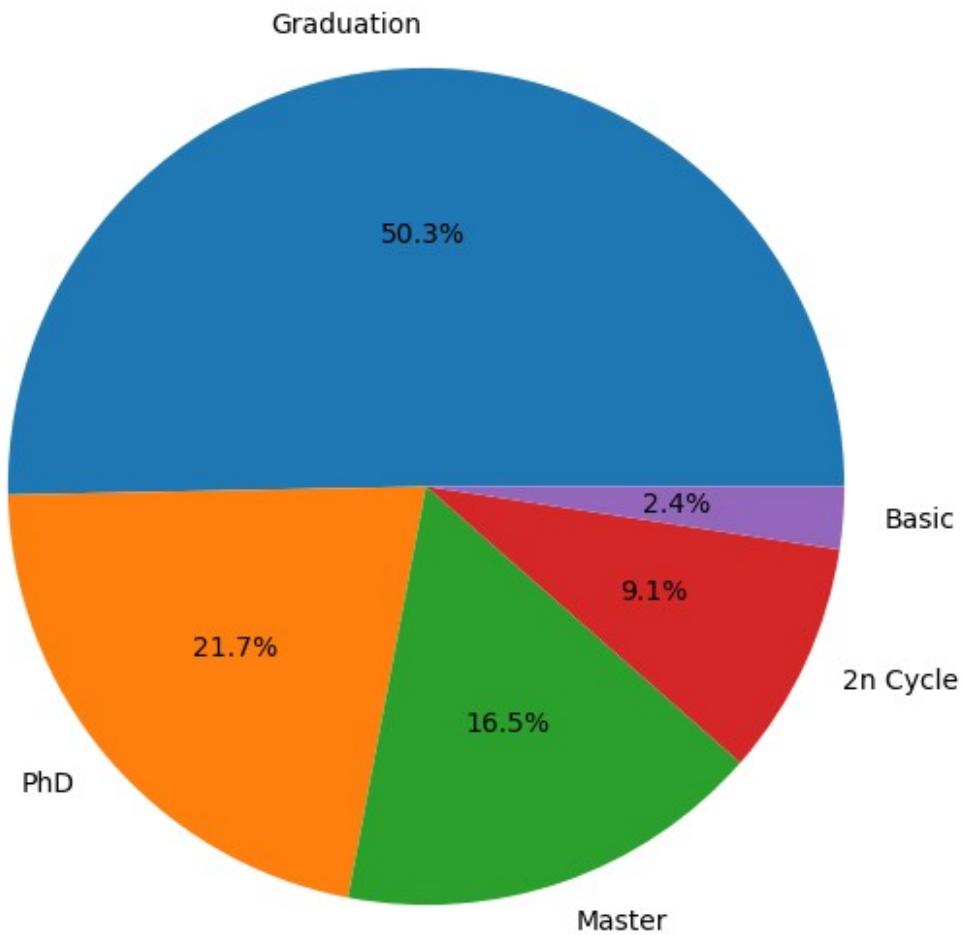
sns.displot(df[ 'Age' ], color='red', kde=True)
plt.show()
```



```
df['Education'].value_counts()

Education
Graduation     1127
PhD            486
Master          370
2n Cycle        203
Basic           54
Name: count, dtype: int64

plt.figure(figsize=(7,7))
ed = df['Education'].value_counts()
plt.pie(ed, autopct='%.1f%'
%) , labels=[ed.index[0],ed.index[1],ed.index[2],ed.index[3],ed.index[4]
])
plt.show()
```



```
# Create a figure with specified size
plt.figure(figsize=(7, 7))

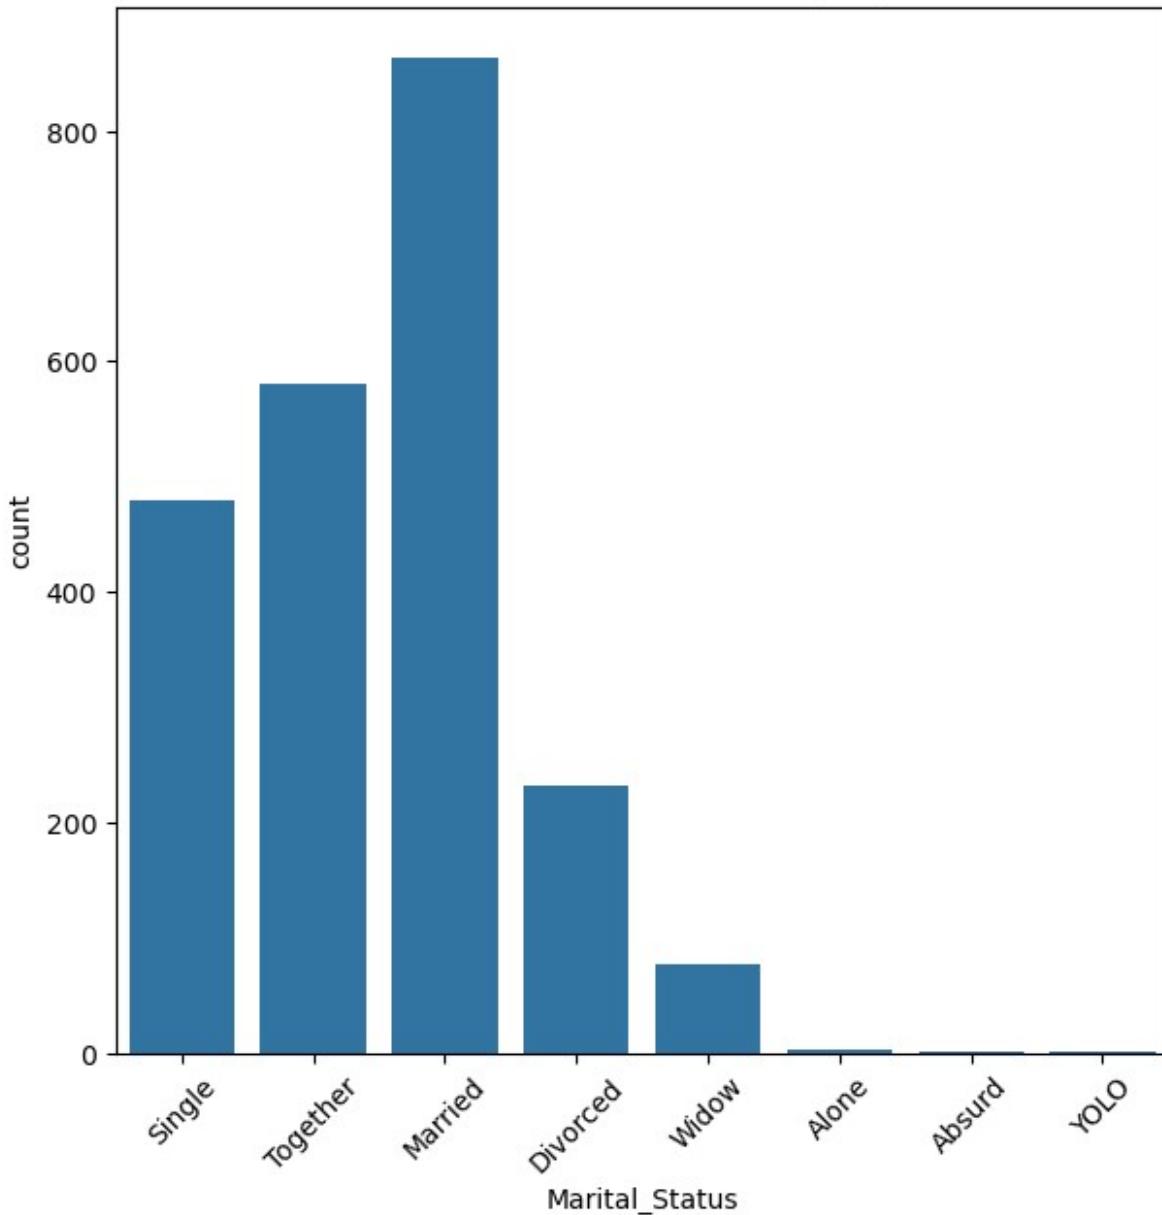
# Create the count plot
ms = sns.countplot(x='Marital_Status', data=df)

# Modify the x-tick labels
ms.tick_params(axis='x', rotation=45)

# Set the title of the plot
plt.title("Count Plot for marital life of people")

# Display the plot
plt.show()
```

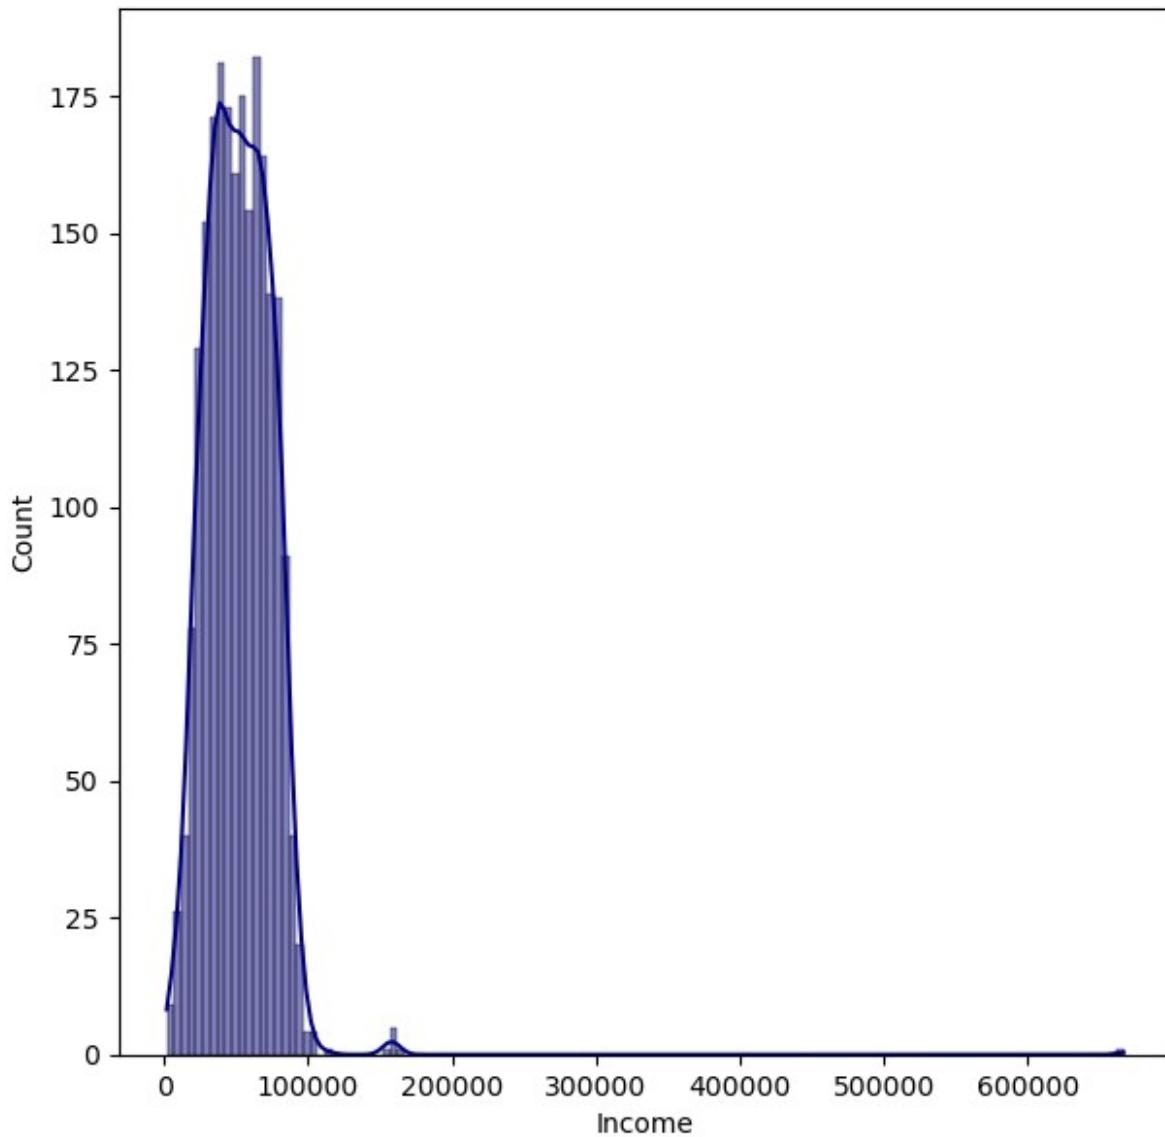
Count Plot for marital life of people



```
# Create a figure with specified size
plt.figure(figsize=(7, 7))

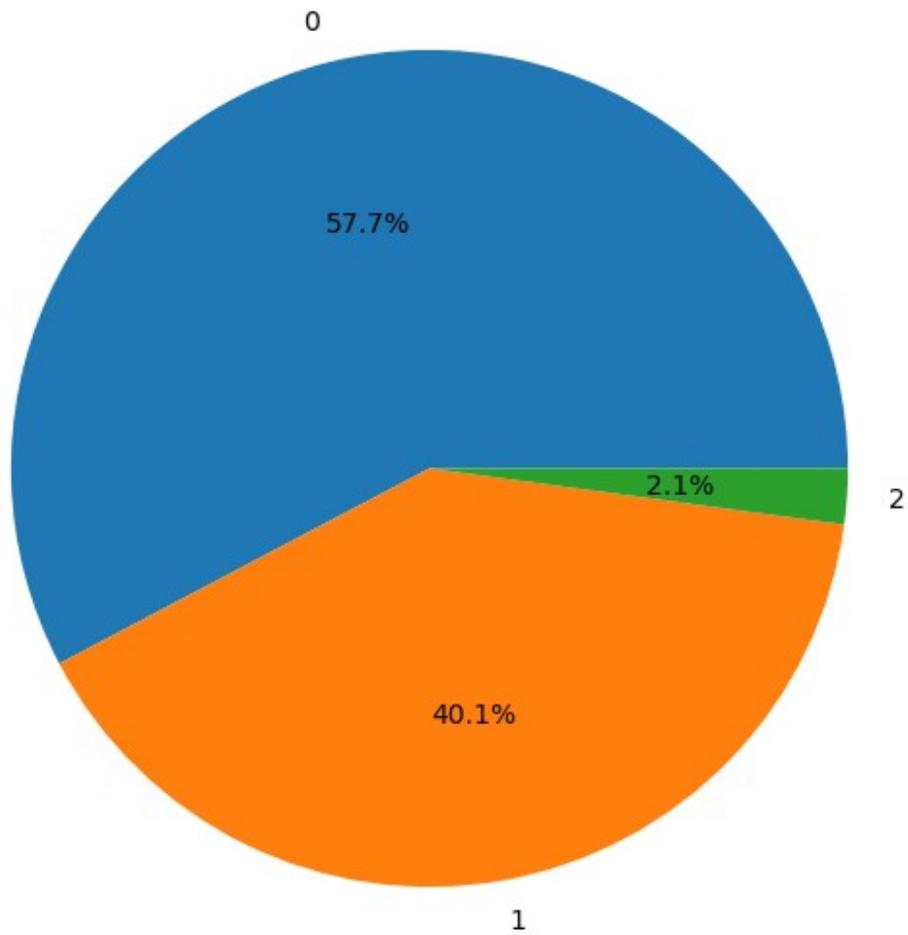
# Create the histogram plot using histplot
sns.histplot(df['Income'], color='navy', kde=True)

# Display the plot
plt.show()
```



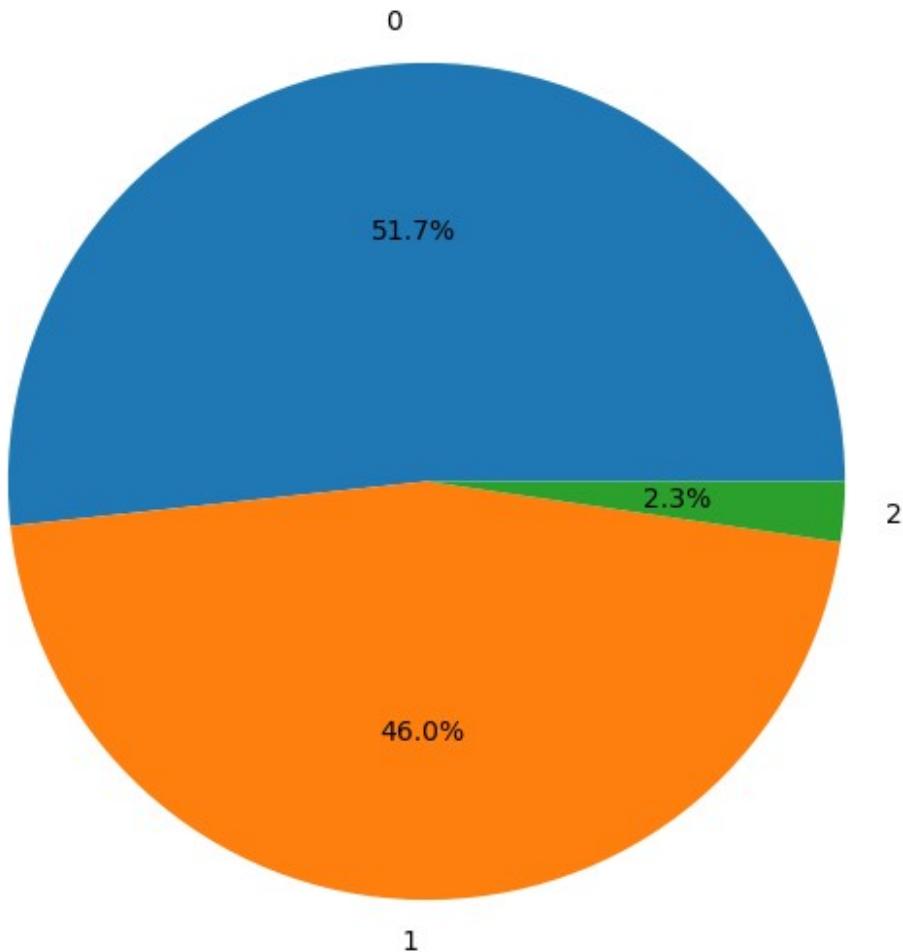
```
plt.figure(figsize=(7,7))
kid = df['Kidhome'].value_counts()
plt.pie(kid, autopct='%.1f%%',
        labels=[kid.index[0], kid.index[1], kid.index[2]])
plt.title("Data for kids available at home")
plt.show()
```

Data for kids available at home



```
plt.figure(figsize=(7,7))
teen= df['Teenhome'].value_counts()
plt.pie(teen, autopct='%.1f%%',
        labels=[teen.index[0],teen.index[1],teen.index[2]])
plt.title("Data for teens available at home")
plt.show()
```

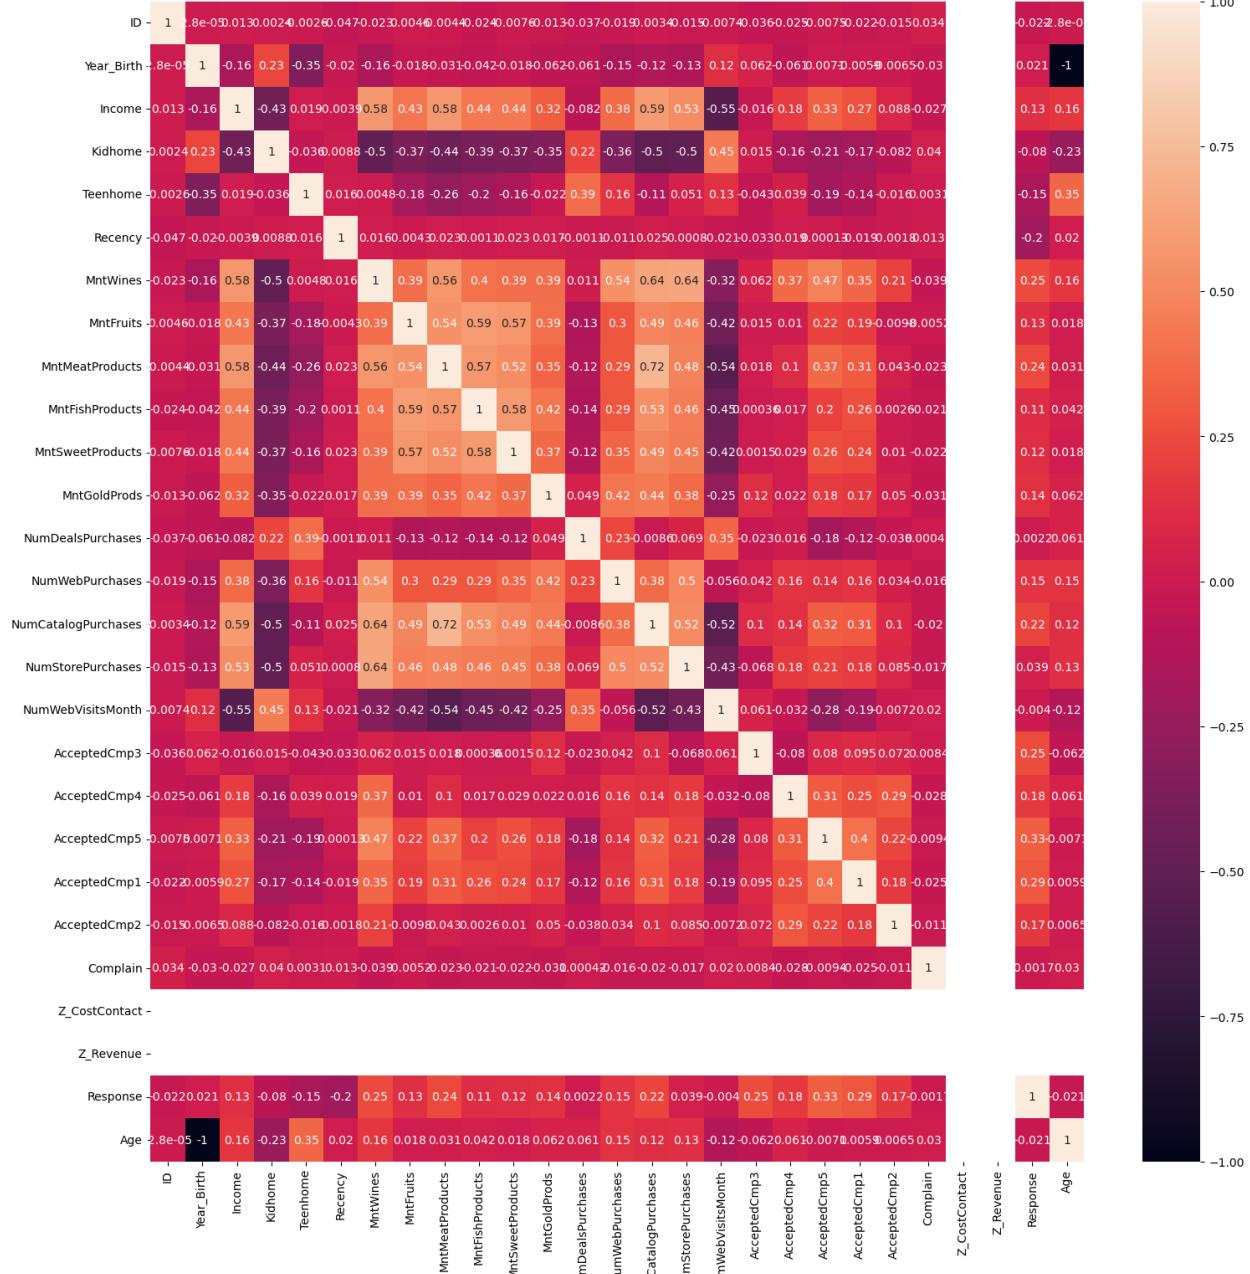
### Data for teens available at home



```
# Select only numeric columns for correlation calculation
numeric_df = df.select_dtypes(include='number')

# Plot the heatmap
plt.figure(figsize=(18, 18))
sns.heatmap(numeric_df.corr(), annot=True)

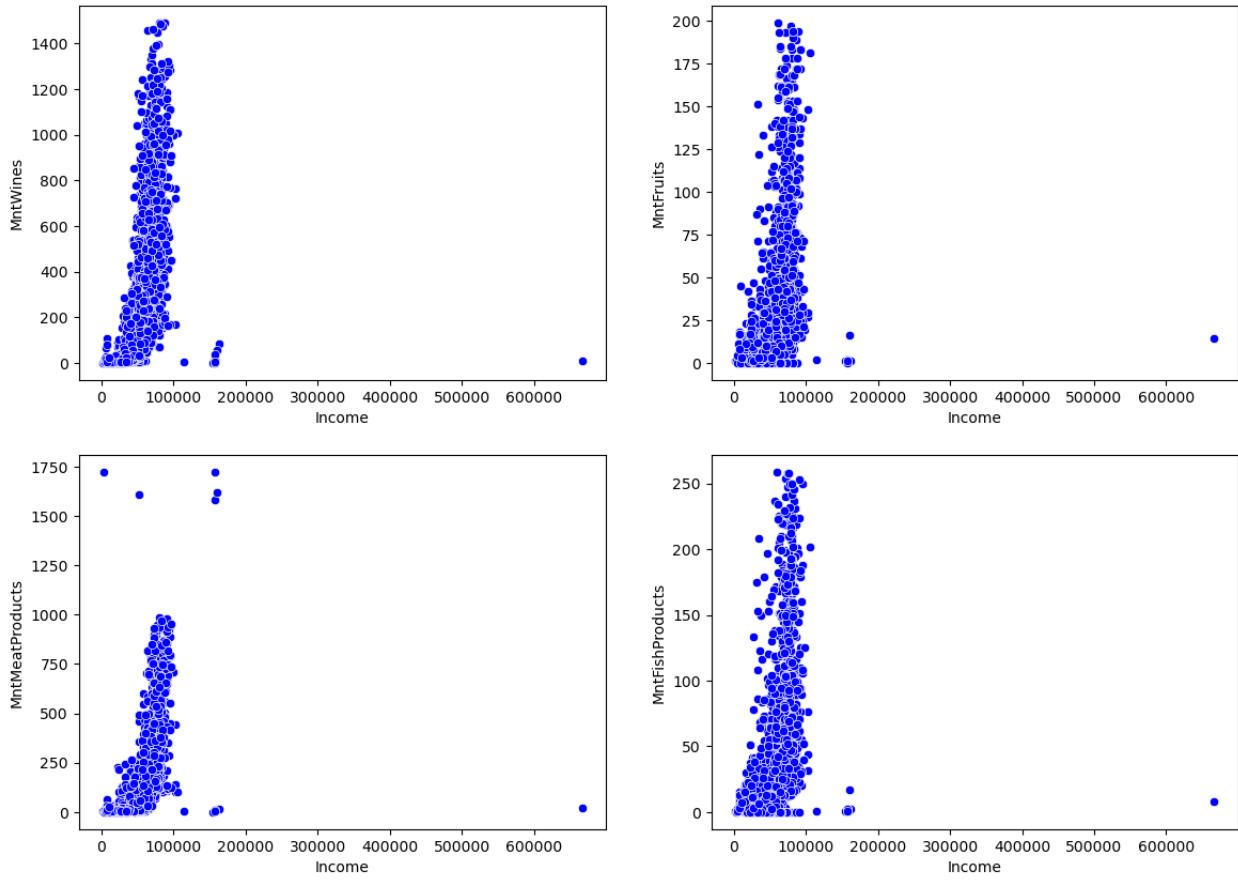
plt.show()
```



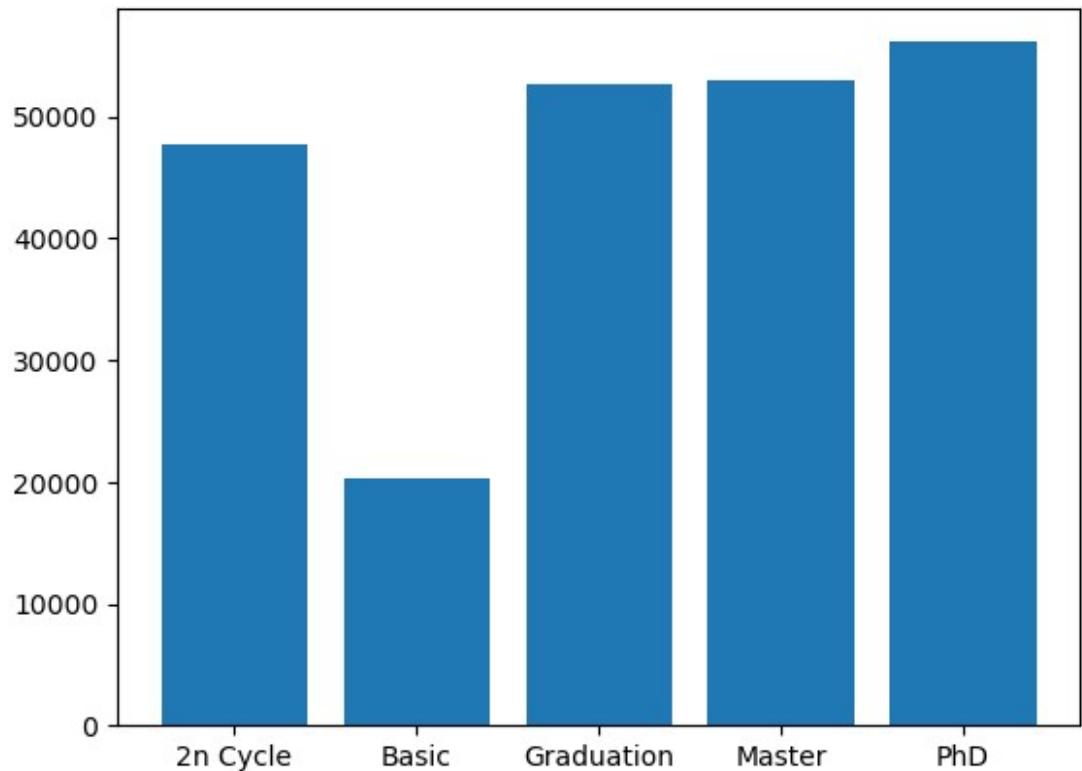
```

plt.figure(figsize=(14,10))
plt.subplot(2,2,1)
sns.scatterplot(data=df,x='Income',y='MntWines',color='blue')
plt.subplot(2,2,2)
sns.scatterplot(data=df,x='Income',y='MntFruits',color='blue')
plt.subplot(2,2,3)
sns.scatterplot(data=df,x='Income',y='MntMeatProducts',color='blue')
plt.subplot(2,2,4)
sns.scatterplot(data=df,x='Income',y='MntFishProducts',color='blue')
plt.show()

```



```
education_income= df.groupby('Education')[['Income']].mean()
plt.bar(education_income.index,height=round(education_income,2))
plt.show()
```



## Data Cleaning and Feature Engineering

```
mean
```

```
52247.25135379061
```

```
df= df.dropna() ## we replaced the income null values with avg or  
mean of income
```

```
df.isnull().sum()
```

ID	0
Year_Birth	0
Education	0
Marital_Status	0
Income	0
Kidhome	0
Teenhome	0
Dt_Customer	0
Recency	0
MntWines	0
MntFruits	0
MntMeatProducts	0
MntFishProducts	0
MntSweetProducts	0
MntGoldProds	0

```

NumDealsPurchases      0
NumWebPurchases        0
NumCatalogPurchases    0
NumStorePurchases      0
NumWebVisitsMonth      0
AcceptedCmp3            0
AcceptedCmp4            0
AcceptedCmp5            0
AcceptedCmp1            0
AcceptedCmp2            0
Complain                0
Z_CostContact          0
Z_Revenue                0
Response                 0
Age                      0
dtype: int64

```

Creating extra features like total money spent, family size

```

df["Total_Spent"] = df["MntWines"]+ df["MntFruits"]+
df["MntMeatProducts"]+ df["MntFishProducts"]+ df["MntSweetProducts"]+
df["MntGoldProds"]

# Replace values in "Marital_Status" column and create "Relation" column
df["Relation"] = df["Marital_Status"].replace({"Married": 2,
"Together": 2, "Absurd": 1, "Widow": 1, "YOLO": 1, "Divorced": 1,
"Single": 1, "Alone": 1}).astype(int)

# Create "Children" column by summing "Kidhome" and "Teenhome"
df["Children"] = df["Kidhome"] + df["Teenhome"]

# Create "Family_Size" column by summing "Relation" and "Children"
df["Family_Size"] = df["Relation"] + df["Children"]

print(df)

```

	ID	Year_Birth	Education	Marital_Status	Income	
Kidhome \	5524	1957	Graduation	Single	58138.0	0
1	2174	1954	Graduation	Single	46344.0	1
2	4141	1965	Graduation	Together	71613.0	0
3	6182	1984	Graduation	Together	26646.0	1
4	5324	1981	PhD	Married	58293.0	1

2235	10870	1967	Graduation	Married	61223.0	0
2236	4001	1946	PhD	Together	64014.0	2
2237	7270	1981	Graduation	Divorced	56981.0	0
2238	8235	1956	Master	Together	69245.0	0
2239	9405	1954	PhD	Married	52869.0	1
Complain \	Teenhome	Dt_Customer	Recency	MntWines	...	AcceptedCmp2
0	0	04-09-2012	58	635	...	0
0	1	08-03-2014	38	11	...	0
0	0	21-08-2013	26	426	...	0
0	0	10-02-2014	26	11	...	0
0	0	19-01-2014	94	173	...	0
0	...	...	...	...	...	...
0	...	...	...	...	...	...
2235	1	13-06-2013	46	709	...	0
0	1	10-06-2014	56	406	...	0
0	0	25-01-2014	91	908	...	0
0	1	24-01-2014	8	428	...	0
0	1	15-10-2012	40	84	...	0
0	...	...	...	...	...	...
Relation \	Z_CostContact	Z_Revenue	Response	Age	Total_Spent	
0	3	11	1	65	1617	1
1	3	11	0	68	27	1
2	3	11	0	57	776	2
3	3	11	0	38	53	2
4	3	11	0	41	422	2

...	...	...	...	...	...	...	...
2235	3	11	0	55	1341	2	
2236	3	11	0	76	444	2	
2237	3	11	0	41	1241	1	
2238	3	11	0	66	843	2	
2239	3	11	1	68	172	2	

	Children	Family_Size
0	0	1
1	2	3
2	0	2
3	1	3
4	1	3
...	...	...
2235	1	3
2236	3	5
2237	0	1
2238	1	3
2239	2	4

[2240 rows x 34 columns]

```
C:\Users\ChatterjeeSo\AppData\Local\Temp\
ipykernel_28452\3846774764.py:2: FutureWarning: Downcasting behavior
in `replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
df["Relation"] = df["Marital_Status"].replace({"Married": 2,
"Together": 2, "Absurd": 1, "Widow": 1, "YOLO": 1, "Divorced": 1,
"Single": 1, "Alone": 1}).astype(int)

df= df.drop(['Relation','Children'],axis=1)

df.columns

Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income',
'Kidhome',
'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits',
'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
'NumCatalogPurchases', 'NumStorePurchases',
'NumWebVisitsMonth',
'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue',
```

```
'Response',
    'Age', 'Total_Spent', 'Family_Size'],
dtype='object')
```

Label Encoding categorical data i.e. Education data

```
from sklearn.preprocessing import LabelEncoder ## one hot encoding
from sklearn.preprocessing import StandardScaler ## scale the values
based on mean of the data

df['Education']

0      Graduation
1      Graduation
2      Graduation
3      Graduation
4          PhD
...
2235    Graduation
2236        PhD
2237    Graduation
2238        Master
2239        PhD
Name: Education, Length: 2240, dtype: object

lb = LabelEncoder()
df['Education']= lb.fit_transform(df['Education'])

df['Response']

0      1
1      0
2      0
3      0
4      0
...
2235    0
2236    0
2237    0
2238    0
2239    1
Name: Response, Length: 2240, dtype: int64
```

Scaling Data

```
df1 = df.copy()
to_drop = ["AcceptedCmp3", "AcceptedCmp4", "AcceptedCmp5",
"AcceptedCmp1","AcceptedCmp2", "Complain", "Response", ]
df1 = df1.drop(to_drop, axis=1)
```

df1

Teenhome \	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
0	5524	1957	2	Single	58138.0	0
1	2174	1954	2	Single	46344.0	1
2	4141	1965	2	Together	71613.0	0
3	6182	1984	2	Together	26646.0	1
4	5324	1981	4	Married	58293.0	1
...	...	...	...	...	...	...
2235	10870	1967	2	Married	61223.0	0
2236	4001	1946	4	Together	64014.0	2
2237	7270	1981	2	Divorced	56981.0	0
2238	8235	1956	3	Together	69245.0	0
2239	9405	1954	4	Married	52869.0	1
NumWebPurchases \	Dt_Customer	Recency	MntWines	...	NumDealsPurchases	
0	04-09-2012	58	635	...	3	
1	08-03-2014	38	11	...	2	
2	21-08-2013	26	426	...	1	
3	10-02-2014	26	11	...	2	
4	19-01-2014	94	173	...	5	
...	...	...	...	...	...	
2235	13-06-2013	46	709	...	2	
2236	10-06-2014	56	406	...	7	
2237	25-01-2014	91	908	...	1	
2238	24-01-2014	8	428	...	2	
2239	15-10-2012	40	84	...	3	

```
3
```

	NumCatalogPurchases	NumStorePurchases	NumWebVisitsMonth	\
0	10	4		7
1	1	2		5
2	2	10		4
3	0	4		6
4	3	6		5
..	..	..	..	..
2235	3	4		5
2236	2	5		7
2237	3	13		6
2238	5	10		3
2239	1	4		7

	Z_CostContact	Z_Revenue	Age	Total_Spent	Family_Size
0	3	11	65	1617	1
1	3	11	68	27	3
2	3	11	57	776	2
3	3	11	38	53	3
4	3	11	41	422	3
..	..	..	..	..	..
2235	3	11	55	1341	3
2236	3	11	76	444	5
2237	3	11	41	1241	1
2238	3	11	66	843	3
2239	3	11	68	172	4

```
[2240 rows x 25 columns]
```

```
scaler = StandardScaler()  
#scaled_feature = scaler.fit_transform(df.values)  
#scaled_df = pd.DataFrame(scaled_feature, index=df.index,  
columns=df.columns)
```

```
df['Response']
```

0	1
1	0
2	0
3	0
4	0
.	.
2235	0
2236	0
2237	0
2238	0
2239	1

```
Name: Response, Length: 2240, dtype: int64
```

# Aproach 1: Unsupervised Learning

## Perform Clustering

We will use the Elbob method to find the optimum number of clusters

```
!pip install yellowbrick
#
Requirement already satisfied: yellowbrick in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (1.5)
Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from yellowbrick) (3.9.2)
Requirement already satisfied: scipy>=1.0.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from yellowbrick) (1.13.1)
Requirement already satisfied: scikit-learn>=1.0.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from yellowbrick) (1.5.1)
Requirement already satisfied: numpy>=1.16.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from yellowbrick) (1.26.4)
Requirement already satisfied: cycler>=0.10.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from yellowbrick) (0.11.0)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.2.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (24.1)
Requirement already satisfied: pillow>=8 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.9.0.post0)
Requirement already satisfied: joblib>=1.2.0 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\
```

```
chatterjeeso\appdata\local\anaconda3\lib\site-packages (from scikit-learn>=1.0.0->yellowbrick) (3.5.0)
Requirement already satisfied: six>=1.5 in c:\users\chatterjeeso\appdata\local\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.16.0)

import os
from sklearn.cluster import KMeans
from sklearn import metrics
from sklearn.cluster import AgglomerativeClustering
from yellowbrick.cluster import KElbowVisualizer

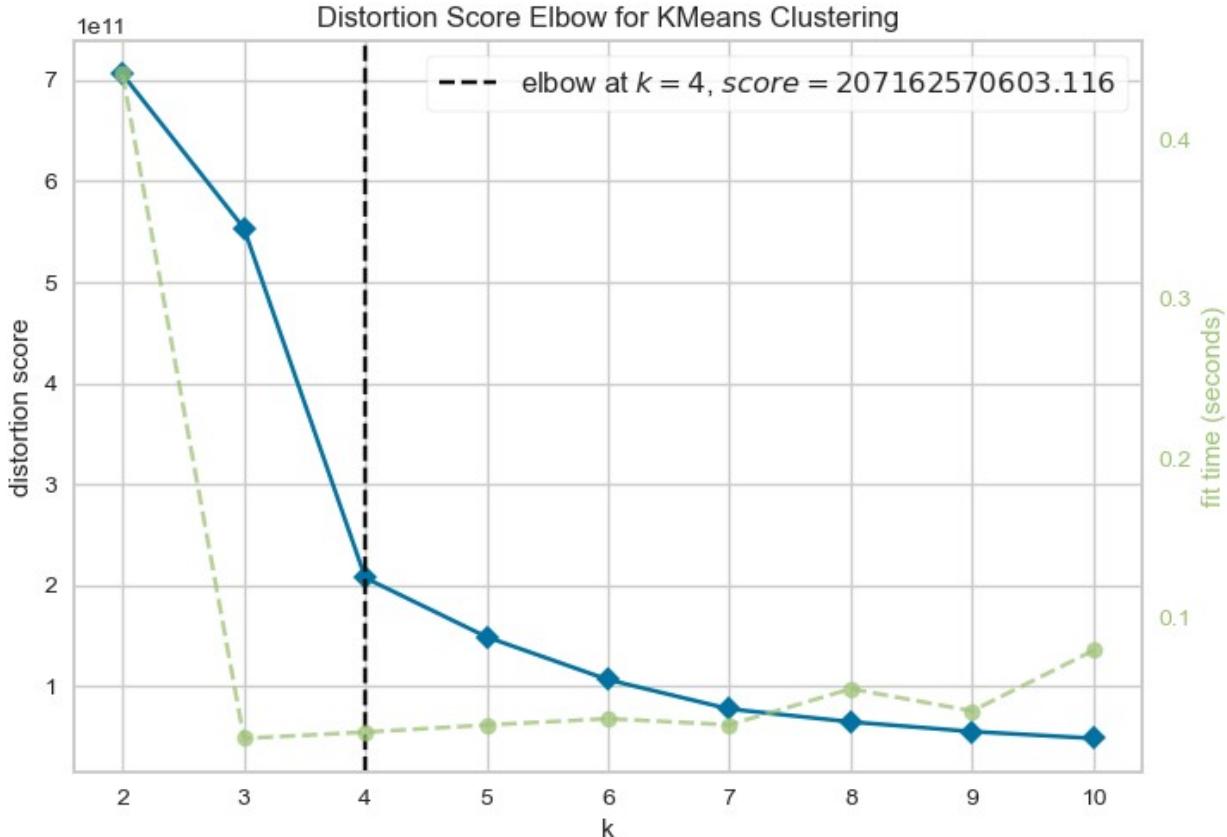
df= df.drop(['Marital_Status','Dt_Customer'],axis=1)

# Set the environment variables to avoid the warnings
os.environ["LOKY_MAX_CPU_COUNT"] = "4"
os.environ["OMP_NUM_THREADS"] = "9"

em = KElbowVisualizer(KMeans(), k=10)
em.fit(df)
em.show()

C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than
```

```
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\
sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have
a memory leak on Windows with MKL, when there are less chunks than
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\
sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have
a memory leak on Windows with MKL, when there are less chunks than
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\
sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have
a memory leak on Windows with MKL, when there are less chunks than
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
    warnings.warn(
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\
sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have
a memory leak on Windows with MKL, when there are less chunks than
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
    warnings.warn(
```



```
<Axes: title={'center': 'Distortion Score Elbow for KMeans Clustering'}, xlabel='k', ylabel='distortion score'>
```

We see that the optimum number of clusters that should be used is k=4

We will use K Means Clustering for the operation

```
from sklearn.cluster import KMeans
os.environ["OMP_NUM_THREADS"] = "9"
kmc = KMeans(n_clusters=4)
# fit model and predict clusters
pred = kmc.fit_predict(df)
df["Clusters"] = pred

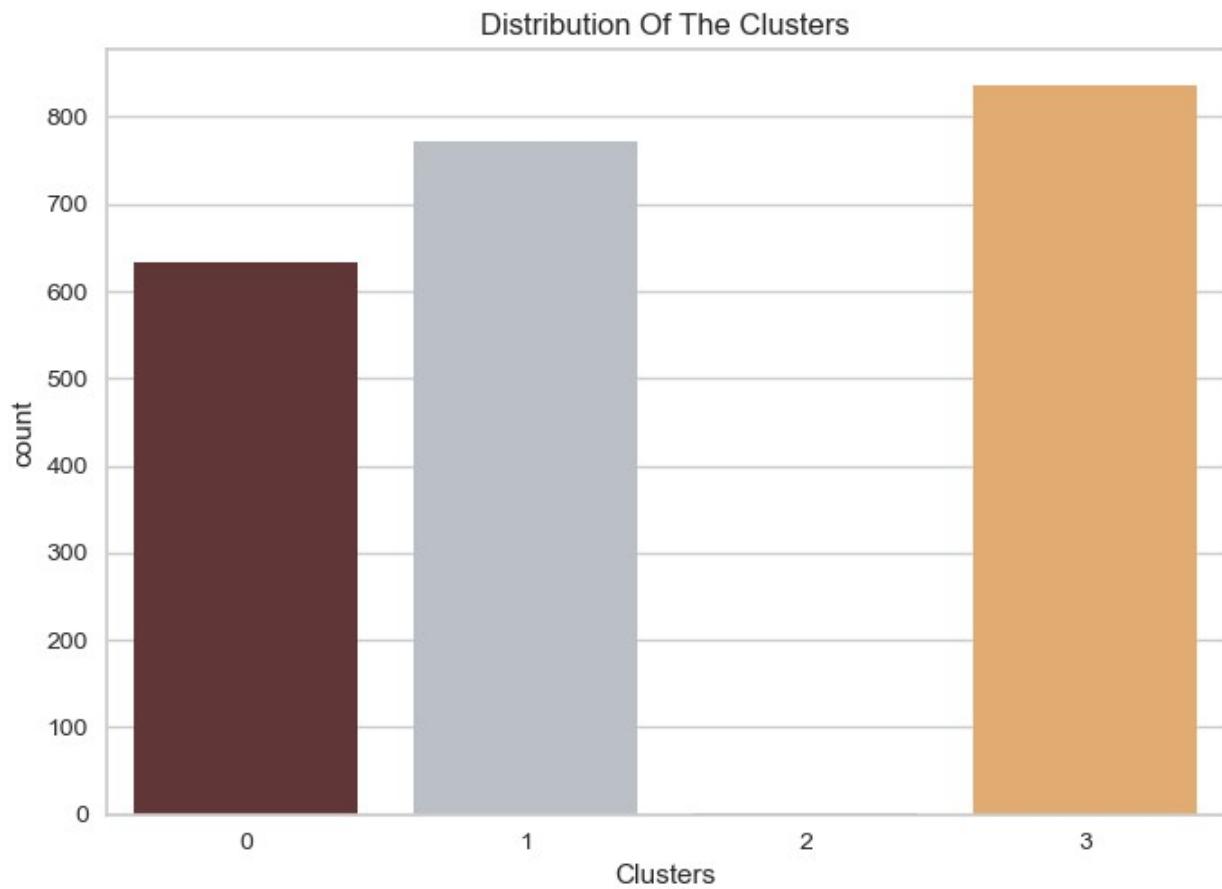
C:\Users\ChatterjeeSo\AppData\Local\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have
a memory leak on Windows with MKL, when there are less chunks than
available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=9.
warnings.warn(
df['Clusters']
```

```
0      3
1      3
2      0
3      1
4      3
 ..
2235    3
2236    3
2237    3
2238    0
2239    3
Name: Clusters, Length: 2240, dtype: int32
```

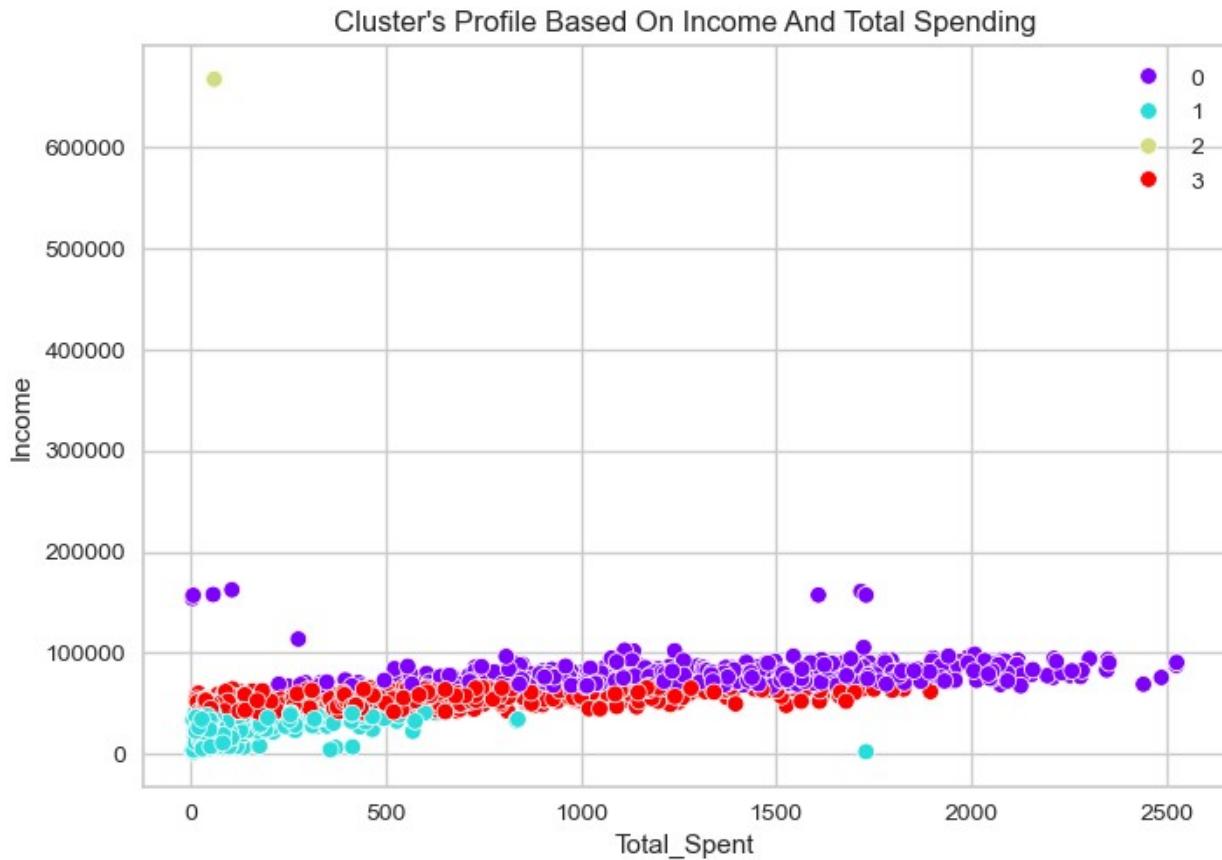
Let us visualize our Clusters

```
# Define the palette
pal = ["#682F2F", "#B9C0C9", "#9F8A78", "#F3AB60"]

# Plot the countplot with hue set to the same variable as x and legend
# set to False
fig = sns.countplot(x="Clusters", hue="Clusters", palette=pal,
                     data=df, legend=False)
fig.set_title("Distribution Of The Clusters")
plt.show()
```



```
fig = sns.scatterplot(data = df,x=df["Total_Spent"],  
y=df["Income"],hue=df["Clusters"], palette="rainbow")  
fig.set_title("Cluster's Profile Based On Income And Total Spending")  
plt.legend()  
plt.show()
```



## Approach 2 : Supervised Learning

```

y=df['Response'] ## dependent var
X_new=df.drop(['Response','Education'],axis=1) ## independent var

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_new,y,
test_size=0.2, random_state=51)

print('Shape of X_train = ', X_train.shape)
print('Shape of y_train = ', y_train.shape)
print('Shape of X_test = ', X_test.shape)
print('Shape of y_test = ', y_test.shape)

Shape of X_train = (1792, 29)
Shape of y_train = (1792,)
Shape of X_test = (448, 29)
Shape of y_test = (448,)
```

## Decision Tree

```

from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier(criterion='gini')
classifier.fit(X_train, y_train)
```

```
DecisionTreeClassifier()  
classifier.score(X_test, y_test)  
0.84375
```

## KNN

```
from sklearn.neighbors import KNeighborsClassifier  
classifier = KNeighborsClassifier(n_neighbors=5)  
classifier.fit(X_train, y_train)  
classifier.score(X_test, y_test)  
  
0.8504464285714286
```

## Random Forest

```
from sklearn.ensemble import RandomForestClassifier  
rf= RandomForestClassifier()  
rf.fit(X_train,y_train)  
RandomForestClassifier()  
rf.score(X_test,y_test)  
  
0.8973214285714286
```

# Customer Lifetime Value (CLV) Prediction Module

```
# --- Customer Lifetime Value (CLV) Prediction Module ---  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Load your data (adjust the path if needed)  
df = pd.read_csv('Customer_Analytics_DB.csv', sep='\t')  
  
# --- Feature Engineering ---  
  
# Fill missing income values with mean (if not already done)  
df['Income'] = df['Income'].fillna(df['Income'].mean())  
  
# Calculate Total_Spent per customer  
product_cols = ['MntWines', 'MntFruits', 'MntMeatProducts',
```

```

'MntFishProducts', 'MntSweetProducts', 'MntGoldProds']
df['Total_Spent'] = df[product_cols].sum(axis=1)

# Calculate Purchase Frequency (total purchases across all channels)
purchase_cols = ['NumWebPurchases', 'NumCatalogPurchases',
                  'NumStorePurchases']
df['Total_Purchases'] = df[purchase_cols].sum(axis=1)

# Estimate Customer Lifespan (in months) using Dt_Customer and Recency
df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'], dayfirst=True,
                                    errors='coerce')
max_date = df['Dt_Customer'].max()
df['Customer_Lifespan_Months'] = ((max_date -
                                    df['Dt_Customer']).dt.days // 30) + 1

# --- CLV Calculation ---
# Simple CLV formula: Average Purchase Value × Purchase Frequency ×
# Lifespan
df['Avg_Purchase_Value'] = df['Total_Spent'] /
    (df['Total_Purchases'].replace(0, np.nan))
df['CLV'] = df['Avg_Purchase_Value'] * df['Total_Purchases'] *
    df['Customer_Lifespan_Months']
df['CLV'] = df['CLV'].fillna(0)

# --- CLV Segmentation ---
# Segment customers into quartiles based on CLV
df['CLV_Segment'] = pd.qcut(df['CLV'], 4, labels=['Low', 'Medium',
                                                 'High', 'Very High'])

# --- Visualization ---
plt.figure(figsize=(10, 6))
sns.histplot(df['CLV'], bins=50, kde=True, color='skyblue')
plt.title('Customer Lifetime Value (CLV) Distribution')
plt.xlabel('CLV')
plt.ylabel('Number of Customers')
plt.show()

plt.figure(figsize=(8, 5))
sns.countplot(x='CLV_Segment', data=df, palette='viridis')
plt.title('CLV Segments')
plt.xlabel('Segment')
plt.ylabel('Number of Customers')
plt.show()

# --- Segment Profile Table ---
segment_profile = df.groupby('CLV_Segment')[['Total_Spent',
                                              'Total_Purchases', 'Income', 'CLV']].mean().round(2)
print("CLV Segment Profiles:\n", segment_profile)

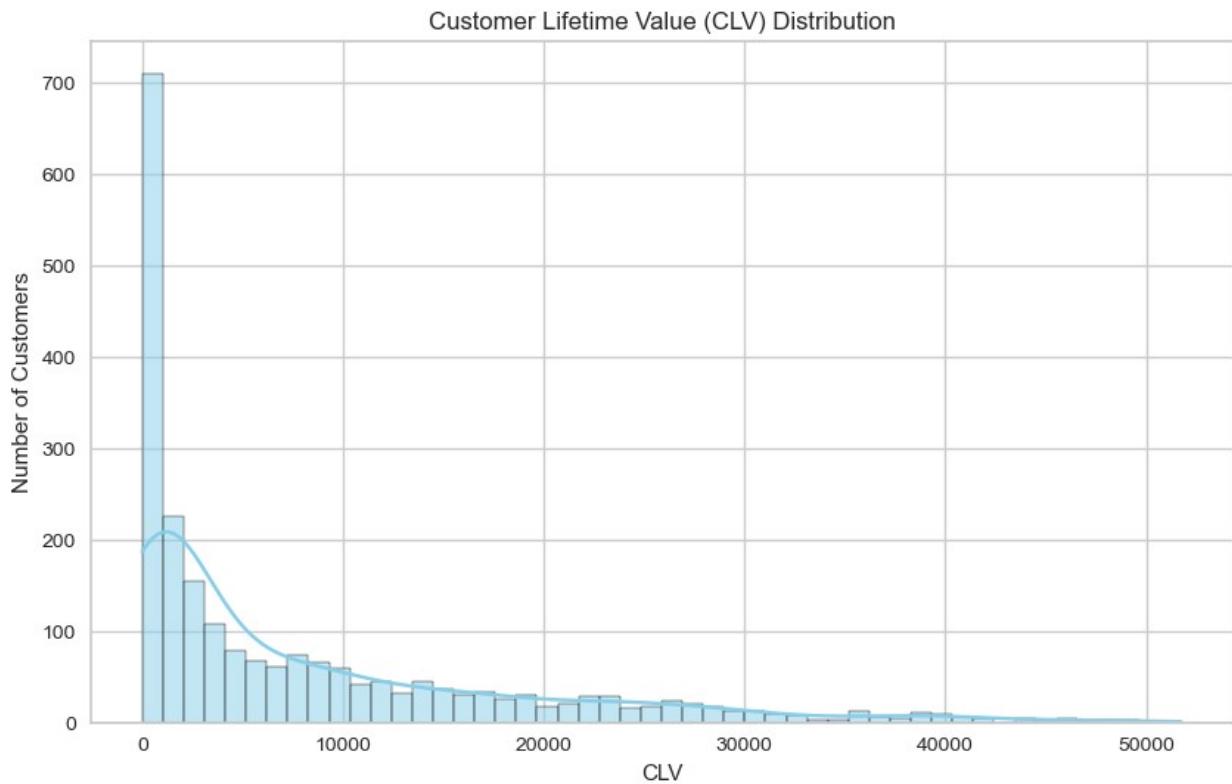
```

```

# --- Insights ---
print("\nTop 5 High-CLV Customers:")
print(df[['ID', 'CLV', 'CLV_Segment', 'Total_Spent',
'Total_Purchases', 'Income']].sort_values('CLV',
ascending=False).head())

# --- Optional: Save results ---
# df.to_csv('customer_clv_results.csv', index=False)

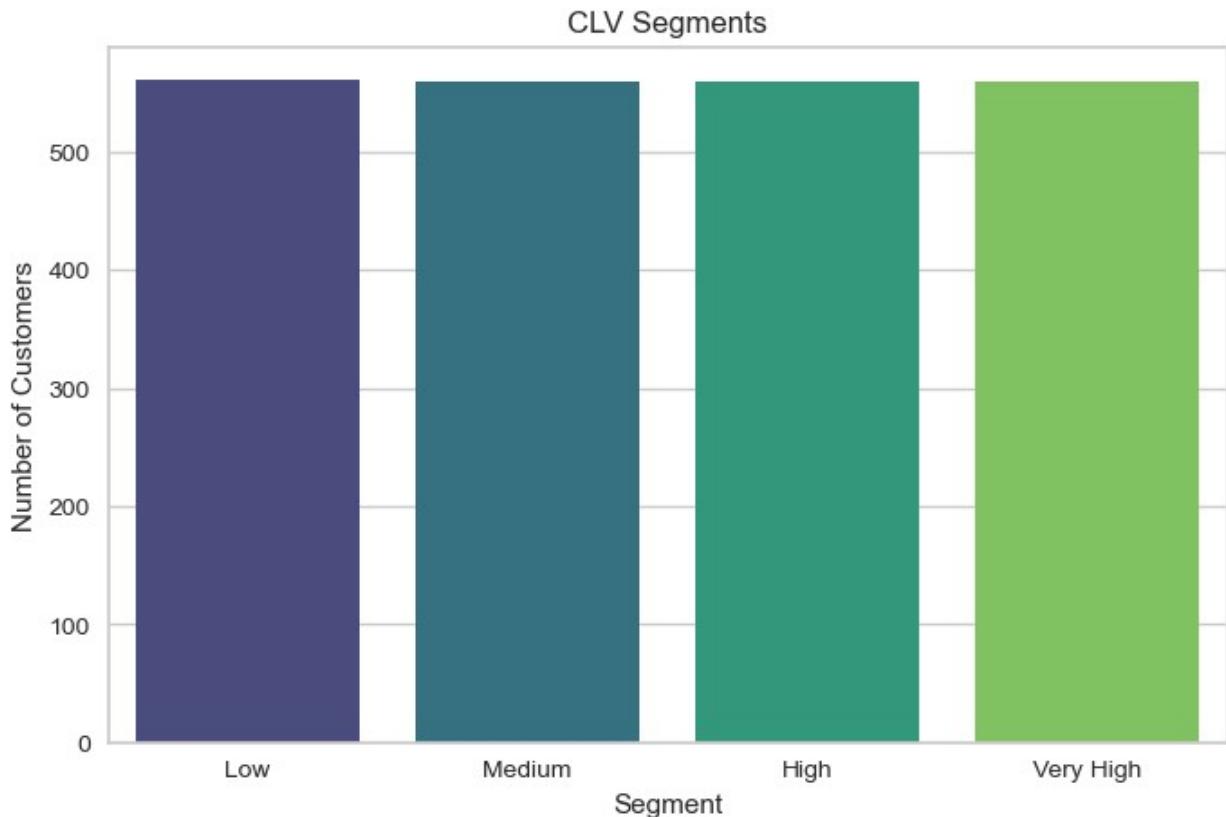
```



```
C:\Users\ChatterjeeSo\AppData\Local\Temp\
ipykernel_28452\1656866214.py:49: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.countplot(x='CLV_Segment', data=df, palette='viridis')
```



#### CLV Segment Profiles:

CLV_Segment	Total_Spent	Total_Purchases	Income	CLV
Low	57.87	4.87	34231.68	298.60
Medium	293.33	9.24	45454.45	1716.93
High	728.90	16.07	58703.37	7183.52
Very High	1343.52	19.97	70619.55	23216.69

#### Top 5 High-CLV Customers:

ID	CLV	CLV_Segment	Total_Spent	Total_Purchases	Income
1601	5453	51744.0	Very High	2352	23 90226.0
1010	5236	50462.0	Very High	2194	20 77568.0
1280	3698	48990.0	Very High	2130	18 78687.0
248	8867	48898.0	Very High	2126	19 67546.0
1890	2747	48898.0	Very High	2126	19 67546.0

```
C:\Users\ChatterjeeSo\AppData\Local\Temp\
ipykernel_28452\1656866214.py:56: FutureWarning: The default of
observed=False is deprecated and will be changed to True in a future
version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
```

```
    segment_profile = df.groupby('CLV_Segment')[['Total_Spent',
'Total_Purchases', 'Income', 'CLV']].mean().round(2)
```